

```
In [52]: import numpy as np
import pandas as pd
import matplotlib.pyplot as py
import seaborn as sns
from sklearn.linear_model import LogisticRegression
```

```
In [53]: df=pd.read_csv(r"D:\New folder\madrid_2003.csv")
df
```

Out[53]:

	date	BEN	CO	EBE	MXV	NMHC	NO_2	NOx	OXY	O_3	PM
0	2003-03-01 01:00:00	NaN	1.72	NaN	NaN	NaN	73.900002	316.299988	NaN	10.550000	55.2099
1	2003-03-01 01:00:00	NaN	1.45	NaN	NaN	0.26	72.110001	250.000000	0.73	6.720000	52.3899
2	2003-03-01 01:00:00	NaN	1.57	NaN	NaN	NaN	80.559998	224.199997	NaN	21.049999	63.2400
3	2003-03-01 01:00:00	NaN	2.45	NaN	NaN	NaN	78.370003	450.399994	NaN	4.220000	67.8399
4	2003-03-01 01:00:00	NaN	3.26	NaN	NaN	NaN	96.250000	479.100006	NaN	8.460000	95.7799
...
243979	2003-10-01 00:00:00	0.20	0.16	2.01	3.17	0.02	31.799999	32.299999	1.68	34.049999	7.3800
243980	2003-10-01 00:00:00	0.32	0.08	0.36	0.72	NaN	10.450000	14.760000	1.00	34.610001	7.4000
243981	2003-10-01 00:00:00	NaN	NaN	NaN	NaN	0.07	34.639999	50.810001	NaN	32.160000	16.8300
243982	2003-10-01 00:00:00	NaN	NaN	NaN	NaN	0.07	32.580002	41.020000	NaN	NaN	13.5700
243983	2003-10-01 00:00:00	1.00	0.29	2.15	6.41	0.07	37.150002	56.849998	2.28	21.480000	12.3500

243984 rows × 16 columns



```
In [3]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 243984 entries, 0 to 243983
Data columns (total 16 columns):
#   Column      Non-Null Count  Dtype
---  -
0   date        243984 non-null object
1   BEN         69745 non-null float64
2   CO          225340 non-null float64
3   EBE         61244 non-null float64
4   MXY         42045 non-null float64
5   NMHC        111951 non-null float64
6   NO_2        242625 non-null float64
7   NOx         242629 non-null float64
8   OXY         42072 non-null float64
9   O_3         234131 non-null float64
10  PM10        240896 non-null float64
11  PXY         42063 non-null float64
12  SO_2        242729 non-null float64
13  TCH         111991 non-null float64
14  TOL         69439 non-null float64
15  station     243984 non-null int64
dtypes: float64(14), int64(1), object(1)
memory usage: 29.8+ MB
```

```
In [4]: df1=df.fillna(value=0)
df1
```

Out[4]:

	date	BEN	CO	EBE	MXV	NMHC	NO_2	NOx	OXY	O_3	PM
0	2003-03-01 01:00:00	0.00	1.72	0.00	0.00	0.00	73.900002	316.299988	0.00	10.550000	55.2099
1	2003-03-01 01:00:00	0.00	1.45	0.00	0.00	0.26	72.110001	250.000000	0.73	6.720000	52.3899
2	2003-03-01 01:00:00	0.00	1.57	0.00	0.00	0.00	80.559998	224.199997	0.00	21.049999	63.2400
3	2003-03-01 01:00:00	0.00	2.45	0.00	0.00	0.00	78.370003	450.399994	0.00	4.220000	67.8399
4	2003-03-01 01:00:00	0.00	3.26	0.00	0.00	0.00	96.250000	479.100006	0.00	8.460000	95.7799
...
243979	2003-10-01 00:00:00	0.20	0.16	2.01	3.17	0.02	31.799999	32.299999	1.68	34.049999	7.3800
243980	2003-10-01 00:00:00	0.32	0.08	0.36	0.72	0.00	10.450000	14.760000	1.00	34.610001	7.4000
243981	2003-10-01 00:00:00	0.00	0.00	0.00	0.00	0.07	34.639999	50.810001	0.00	32.160000	16.8300
243982	2003-10-01 00:00:00	0.00	0.00	0.00	0.00	0.07	32.580002	41.020000	0.00	0.000000	13.5700
243983	2003-10-01 00:00:00	1.00	0.29	2.15	6.41	0.07	37.150002	56.849998	2.28	21.480000	12.3500

243984 rows × 16 columns



```
In [5]: df1.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 243984 entries, 0 to 243983
Data columns (total 16 columns):
#   Column      Non-Null Count  Dtype
---  -
0   date        243984 non-null  object
1   BEN         243984 non-null  float64
2   CO          243984 non-null  float64
3   EBE         243984 non-null  float64
4   MXY         243984 non-null  float64
5   NMHC        243984 non-null  float64
6   NO_2        243984 non-null  float64
7   NOx         243984 non-null  float64
8   OXY         243984 non-null  float64
9   O_3         243984 non-null  float64
10  PM10        243984 non-null  float64
11  PXY         243984 non-null  float64
12  SO_2        243984 non-null  float64
13  TCH         243984 non-null  float64
14  TOL         243984 non-null  float64
15  station     243984 non-null  int64
dtypes: float64(14), int64(1), object(1)
memory usage: 29.8+ MB
```

```
In [6]: df1.columns
```

```
Out[6]: Index(['date', 'BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
               'PM10', 'PXY', 'SO_2', 'TCH', 'TOL', 'station'],
              dtype='object')
```

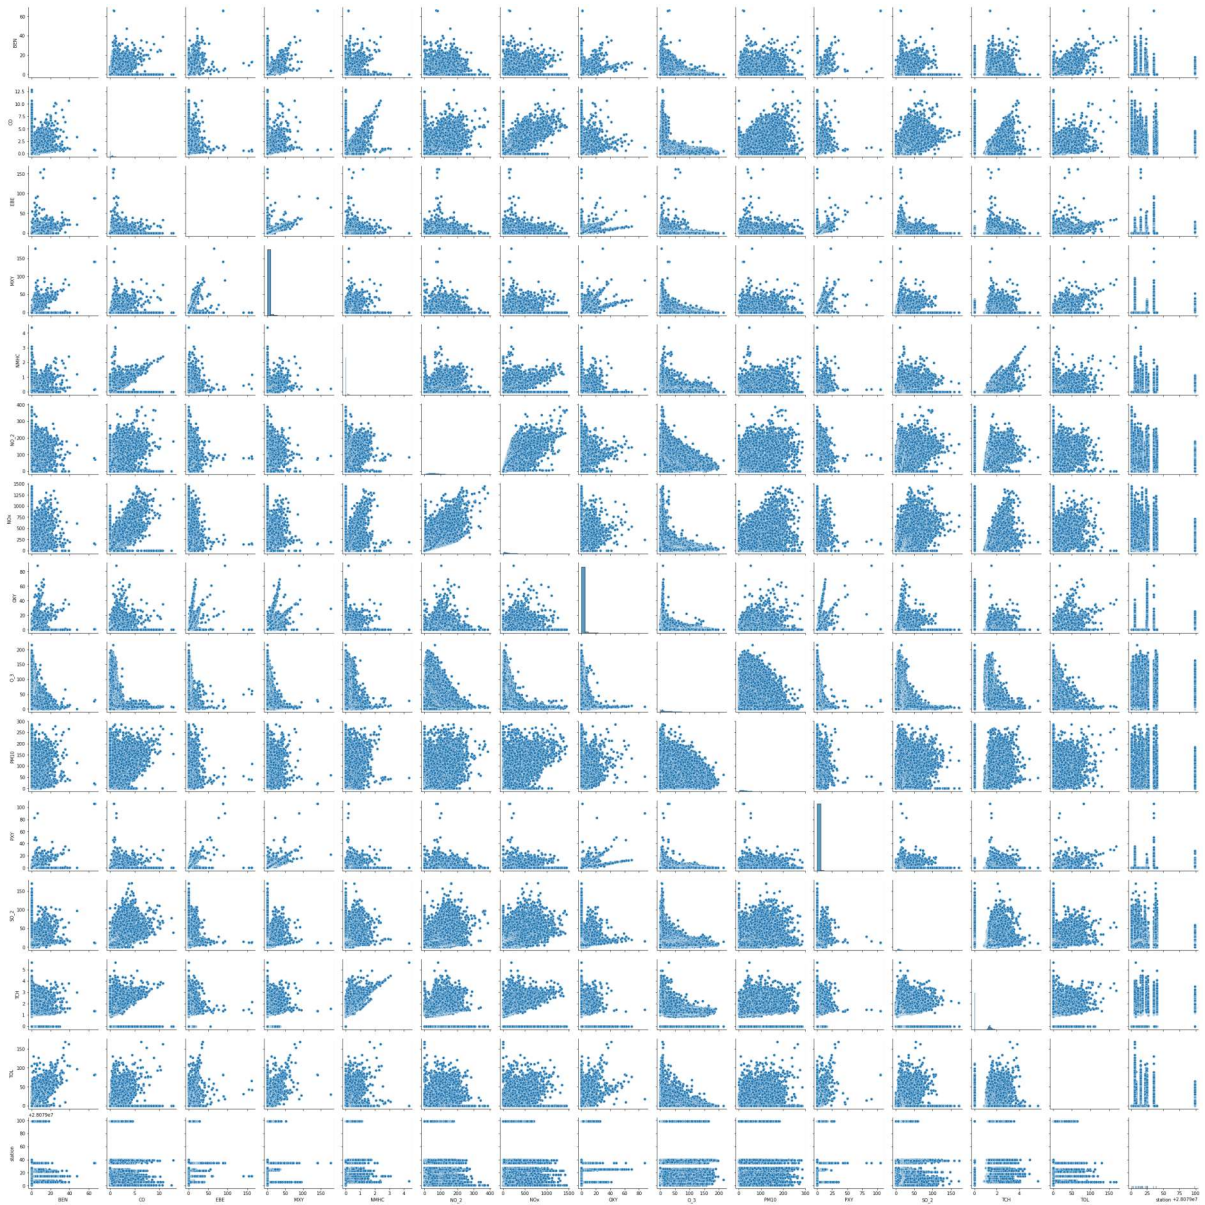
```
In [7]: df2=df1[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',  
              'PM10', 'PXY', 'SO_2', 'TCH', 'TOL', 'station']]  
df2
```

Out[7]:

	BEN	CO	EBE	MXY	NMHC	NO_2	NOx	OXY	O_3	PM10	PXY
0	0.00	1.72	0.00	0.00	0.00	73.900002	316.299988	0.00	10.550000	55.209999	0.0
1	0.00	1.45	0.00	0.00	0.26	72.110001	250.000000	0.73	6.720000	52.389999	0.0
2	0.00	1.57	0.00	0.00	0.00	80.559998	224.199997	0.00	21.049999	63.240002	0.0
3	0.00	2.45	0.00	0.00	0.00	78.370003	450.399994	0.00	4.220000	67.839996	0.0
4	0.00	3.26	0.00	0.00	0.00	96.250000	479.100006	0.00	8.460000	95.779999	0.0
...
243979	0.20	0.16	2.01	3.17	0.02	31.799999	32.299999	1.68	34.049999	7.380000	1.2
243980	0.32	0.08	0.36	0.72	0.00	10.450000	14.760000	1.00	34.610001	7.400000	0.5
243981	0.00	0.00	0.00	0.00	0.07	34.639999	50.810001	0.00	32.160000	16.830000	0.0
243982	0.00	0.00	0.00	0.00	0.07	32.580002	41.020000	0.00	0.000000	13.570000	0.0
243983	1.00	0.29	2.15	6.41	0.07	37.150002	56.849998	2.28	21.480000	12.350000	2.4

```
In [8]: sns.pairplot(df2)
```

```
Out[8]: <seaborn.axisgrid.PairGrid at 0x26b0892beb0>
```




```
In [9]: sns.distplot(df2['station'])
```

```
C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
```

```
warnings.warn(msg, FutureWarning)
```

```
Out[9]: <AxesSubplot:xlabel='station', ylabel='Density'>
```



```
In [10]: x=df2[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'O_3', 'PM10', 'PXY',
y=df2['station']]
```

```
In [11]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.3)
```

linear

```
In [12]: from sklearn.linear_model import LinearRegression
```

```
In [13]: lr=LinearRegression()
         lr.fit(x_train,y_train)
```

Out[13]: LinearRegression()

```
In [14]: coeff =pd.DataFrame(lr.coef_,x.columns,columns=["Co-efficient"])
coeff
```

Out[14]:


	Co-efficient
BEN	1.329314
CO	0.078055
EBE	-1.103965
MXY	0.395101
NMHC	0.744723
NO_2	-0.079070
NOx	-0.008685
O_3	-0.019030
PM10	0.043117
PXY	2.733720
SO_2	-0.156103
TCH	4.654519
TOL	-0.231877
OXY	0.262578

```
In [15]: print(lr.intercept_)
```

28079024.492372576

```
In [16]: prediction =lr.predict(x_test)
py.scatter(y_test,prediction)
```

Out[16]: <matplotlib.collections.PathCollection at 0x26b318a5280>




```
In [17]: print(lr.score(x_test,y_test))
```

```
0.12427344825624764
```

```
In [18]: print(lr.score(x_train,y_train))
```

```
0.11829138820457574
```

Ridge

```
In [19]: from sklearn.linear_model import Ridge,Lasso
```

```
In [20]: rr=Ridge(alpha=10)  
rr.fit(x_train,y_train)
```

```
Out[20]: Ridge(alpha=10)
```

```
In [21]: rr.score(x_test,y_test)
```

```
Out[21]: 0.1242719430594611
```

Lasso

```
In [22]: la=Lasso(alpha=10)  
la.fit(x_train,y_train)
```

```
Out[22]: Lasso(alpha=10)
```

```
In [23]: la.score(x_test,y_test)
```

```
Out[23]: 0.040051613031118816
```

elasticnet

```
In [24]: from sklearn.linear_model import ElasticNet  
en=ElasticNet()  
en.fit(x_train,y_train)
```

```
Out[24]: ElasticNet()
```

```
In [25]: print(en.coef_)
```

```
[ 0.13012942 -0.          -0.          0.89170003  0.          -0.07099582  
 -0.00816406 -0.01816097  0.04245185  0.56079423 -0.17328773  1.78537876  
  0.          0.17031933]
```

```
In [26]: print(en.intercept_)
```

```
28079025.995923337
```

```
In [27]: print(en.predict(x_test))
```

```
[28079027.46757036 28079021.1167562 28079022.89446069 ...  
28079025.9103636 28079025.96325697 28079027.48770221]
```

```
In [28]: print(en.score(x_test,y_test))
```

```
0.10087740966978953
```

logistic

```
In [29]: feature_matrix=df2.iloc[:,0:14]  
target_vector=df2.iloc[:,-1]
```

```
In [30]: feature_matrix=df2[['BEN', 'CO', 'EBE', 'MXV', 'NMHC', 'NO_2', 'NOx', 'O_3', 'PM  
y=df2['station']
```

```
In [31]: feature_matrix.shape
```

```
Out[31]: (243984, 14)
```

```
In [32]: target_vector.shape
```

```
Out[32]: (243984,)
```

```
In [33]: from sklearn.preprocessing import StandardScaler
```

```
In [34]: fs=StandardScaler().fit_transform(feature_matrix)
```

```
In [35]: logr =LogisticRegression()  
logr.fit(fs,target_vector)
```

```
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\_logistic.py:  
763: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

```
Increase the number of iterations (max_iter) or scale the data as shown in:  
https://scikit-learn.org/stable/modules/preprocessing.html (https://scikit-learn.org/stable/modules/preprocessing.html)  
Please also refer to the documentation for alternative solver options:  
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)  
n_iter_i = _check_optimize_result(  

```

```
Out[35]: LogisticRegression()
```

```
In [36]: observation=[[1.4,2.3,5.0,11,12,13,14,15,4,5,7,6,7,13]]
```

```
In [37]: prediction=logr.predict(observation)
print(prediction)

[28079099]
```

```
In [38]: logr.classes_
```

```
Out[38]: array([28079001, 28079003, 28079004, 28079006, 28079007, 28079008,
                28079009, 28079011, 28079012, 28079014, 28079015, 28079016,
                28079017, 28079018, 28079019, 28079021, 28079022, 28079023,
                28079024, 28079025, 28079026, 28079027, 28079035, 28079036,
                28079038, 28079039, 28079040, 28079099], dtype=int64)
```

```
In [39]: logr.score(fs,target_vector)
```

```
Out[39]: 0.9353441209259623
```

```
In [40]: logr.predict_proba(observation)[0][0]
```

```
Out[40]: 0.0
```

```
In [41]: logr.predict_proba(observation)[0][1]
```

```
Out[41]: 0.0
```

random forest

```
In [42]: from sklearn.ensemble import RandomForestClassifier
from sklearn.tree import plot_tree
```

```
In [43]: x=df2.drop('station',axis=1)
y=df2['station']
```

```
In [44]: x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.70)
```

```
In [45]: rfc=RandomForestClassifier()
rfc.fit(x_train,y_train)
```

```
Out[45]: RandomForestClassifier()
```

```
In [46]: parameters = {'max_depth':[1,2,3,4,5],
                        'min_samples_leaf':[5,10,15,20,25],
                        'n_estimators':[10,20,30,40,50]}
```

```
In [47]: from sklearn.model_selection import GridSearchCV
```

