In [2]:
```python
import numpy as np
import pandas as pd
import matplotlib.pyplot as py
import seaborn as sns
from sklearn.linear_model import LogisticRegression
```

In [3]:
```python
df=pd.read_csv(r"D:\New folder\madrid_2013.csv")
df
```

Out[3]:

	date	BEN	CO	EBE	MXY	NMHC	NO_2	NOx	OXY	O_3	PN
0	2003-03-01 01:00:00	NaN	1.72	NaN	NaN	NaN	73.900002	316.299988	NaN	10.550000	55.2099
1	2003-03-01 01:00:00	NaN	1.45	NaN	NaN	0.26	72.110001	250.000000	0.73	6.720000	52.3899
2	2003-03-01 01:00:00	NaN	1.57	NaN	NaN	NaN	80.559998	224.199997	NaN	21.049999	63.2400
3	2003-03-01 01:00:00	NaN	2.45	NaN	NaN	NaN	78.370003	450.399994	NaN	4.220000	67.8399
4	2003-03-01 01:00:00	NaN	3.26	NaN	NaN	NaN	96.250000	479.100006	NaN	8.460000	95.7799
...	...	...	...	...	...	...	...	...	...	...	
243979	2003-10-01 00:00:00	0.20	0.16	2.01	3.17	0.02	31.799999	32.299999	1.68	34.049999	7.3800
243980	2003-10-01 00:00:00	0.32	0.08	0.36	0.72	NaN	10.450000	14.760000	1.00	34.610001	7.4000
243981	2003-10-01 00:00:00	NaN	NaN	NaN	NaN	0.07	34.639999	50.810001	NaN	32.160000	16.8300
243982	2003-10-01 00:00:00	NaN	NaN	NaN	NaN	0.07	32.580002	41.020000	NaN	NaN	13.5700
243983	2003-10-01 00:00:00	1.00	0.29	2.15	6.41	0.07	37.150002	56.849998	2.28	21.480000	12.3500

243984 rows × 16 columns
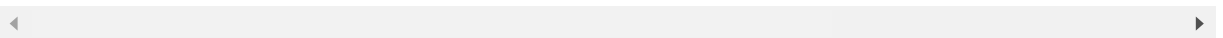
In [3]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 217872 entries, 0 to 217871
Data columns (total 16 columns):
 #   Column   Non-Null Count   Dtype
---  ------   --------------   -----
 0   date     217872 non-null  object
 1   BEN      70389 non-null   float64
 2   CO       216341 non-null  float64
 3   EBE      57752 non-null   float64
 4   MXY      42753 non-null   float64
 5   NMHC     85719 non-null   float64
 6   NO_2     216331 non-null  float64
 7   NOx      216318 non-null  float64
 8   OXY      42856 non-null   float64
 9   O_3      216514 non-null  float64
 10  PM10     207776 non-null  float64
 11  PXY      42845 non-null   float64
 12  SO_2     216403 non-null  float64
 13  TCH      85797 non-null   float64
 14  TOL      70196 non-null   float64
 15  station  217872 non-null  int64
dtypes: float64(14), int64(1), object(1)
memory usage: 26.6+ MB
```

In [4]:
```python
df1=df.fillna(value=0)
df1
```

Out[4]:

	date	BEN	CO	EBE	MXY	NMHC	NO_2	NOx	OXY	O_3	
0	2001-08-01 01:00:00	0.00	0.37	0.00	0.00	0.00	58.400002	87.150002	0.00	34.529999	105.0(
1	2001-08-01 01:00:00	1.50	0.34	1.49	4.10	0.07	56.250000	75.169998	2.11	42.160000	100.5!
2	2001-08-01 01:00:00	0.00	0.28	0.00	0.00	0.00	50.660000	61.380001	0.00	46.310001	100.0!
3	2001-08-01 01:00:00	0.00	0.47	0.00	0.00	0.00	69.790001	73.449997	0.00	40.650002	69.7
4	2001-08-01 01:00:00	0.00	0.39	0.00	0.00	0.00	22.830000	24.799999	0.00	66.309998	75.1
...	...	...	...	...	...	...	...	...	...	...	
217867	2001-04-01 00:00:00	10.45	1.81	0.00	0.00	0.00	73.000000	264.399994	0.00	5.200000	47.8
217868	2001-04-01 00:00:00	5.20	0.69	4.56	0.00	0.13	71.080002	129.300003	0.00	13.460000	26.8(
217869	2001-04-01 00:00:00	0.49	1.09	0.00	1.00	0.19	76.279999	128.399994	0.35	5.020000	40.7
217870	2001-04-01 00:00:00	5.62	1.01	5.04	11.38	0.00	80.019997	197.000000	2.58	5.840000	37.8
217871	2001-04-01 00:00:00	8.09	1.62	6.66	13.04	0.18	76.809998	206.300003	5.20	8.340000	35.3(

217872 rows × 16 columns

In [5]: `df1.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 217872 entries, 0 to 217871
Data columns (total 16 columns):
 #   Column   Non-Null Count   Dtype
---  ------   --------------   -----
 0   date     217872 non-null  object
 1   BEN      217872 non-null  float64
 2   CO       217872 non-null  float64
 3   EBE      217872 non-null  float64
 4   MXY      217872 non-null  float64
 5   NMHC     217872 non-null  float64
 6   NO_2     217872 non-null  float64
 7   NOx      217872 non-null  float64
 8   OXY      217872 non-null  float64
 9   O_3      217872 non-null  float64
 10  PM10     217872 non-null  float64
 11  PXY      217872 non-null  float64
 12  SO_2     217872 non-null  float64
 13  TCH      217872 non-null  float64
 14  TOL      217872 non-null  float64
 15  station  217872 non-null  int64
dtypes: float64(14), int64(1), object(1)
memory usage: 26.6+ MB
```

In [6]: `df1.columns`

Out[6]: `Index(['date', 'BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',`
        `       'PM10', 'PXY', 'SO_2', 'TCH', 'TOL', 'station'],`
        `      dtype='object')`

In [7]:
```python
df2=df1[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
         'PM10', 'PXY', 'SO_2', 'TCH', 'TOL', 'station']]
df2
```

Out[7]:

	BEN	CO	EBE	MXY	NMHC	NO_2	NOx	OXY	O_3	PM10
0	0.00	0.37	0.00	0.00	0.00	58.400002	87.150002	0.00	34.529999	105.000000
1	1.50	0.34	1.49	4.10	0.07	56.250000	75.169998	2.11	42.160000	100.599998
2	0.00	0.28	0.00	0.00	0.00	50.660000	61.380001	0.00	46.310001	100.099998
3	0.00	0.47	0.00	0.00	0.00	69.790001	73.449997	0.00	40.650002	69.779999
4	0.00	0.39	0.00	0.00	0.00	22.830000	24.799999	0.00	66.309998	75.180000
...	...	...	...	...	...	...	...	...	...	...
217867	10.45	1.81	0.00	0.00	0.00	73.000000	264.399994	0.00	5.200000	47.880001
217868	5.20	0.69	4.56	0.00	0.13	71.080002	129.300003	0.00	13.460000	26.809999
217869	0.49	1.09	0.00	1.00	0.19	76.279999	128.399994	0.35	5.020000	40.770000
217870	5.62	1.01	5.04	11.38	0.00	80.019997	197.000000	2.58	5.840000	37.889999
217871	8.09	1.62	6.66	13.04	0.18	76.809998	206.300003	5.20	8.340000	35.369999

In [ ]:
```python
sns.pairplot(df2)
```

Out[8]:
```
<seaborn.axisgrid.PairGrid at 0x2e2bb15a040>
```

In [ ]:
```python
sns.distplot(df2['station'])
```

In [ ]:
```python
x=df2[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'O_3','PM10', 'PXY',
y=df2['station']
```

In [ ]:
```python
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.3)
```

linear

In [ ]:
```python
from sklearn.linear_model import LinearRegression
```

In [ ]:
```python
lr=LinearRegression()
lr.fit(x_train,y_train)
```

In [ ]:
```python
coeff =pd.DataFrame(lr.coef_,x.columns,columns=["Co-efficient"])
coeff
```

In [ ]:
```python
print(lr.intercept_)
```

```
In [ ]: prediction =lr.predict(x_test)
        py.scatter(y_test,prediction)
```

```
In [ ]: print(lr.score(x_test,y_test))
```

```
In [ ]: print(lr.score(x_train,y_train))
```

Ridge

```
In [ ]: from sklearn.linear_model import Ridge,Lasso
```

```
In [ ]: rr=Ridge(alpha=10)
        rr.fit(x_train,y_train)
```

```
In [ ]: rr.score(x_test,y_test)
```

Lasso

```
In [ ]: la=Lasso(alpha=10)
        la.fit(x_train,y_train)
```

```
In [ ]: la.score(x_test,y_test)
```

elasticnet

```
In [ ]: from sklearn.linear_model import ElasticNet
        en=ElasticNet()
        en.fit(x_train,y_train)
```

```
In [ ]: print(en.coef_)
```

```
In [ ]: print(en.intercept_)
```

```
In [ ]: print(en.predict(x_test))
```

```
In [ ]: print(en.score(x_test,y_test))
```

logistic

```python
feature_matrix=df2.iloc[:,0:14]
target_vector=df2.iloc[:,-1]
```

```python
feature_matrix=df2[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'O_3','PI
y=df2['station']
```

```python
feature_matrix.shape
```

```python
target_vector.shape
```

```python
from sklearn.preprocessing import StandardScaler
```

```python
fs=StandardScaler().fit_transform(feature_matrix)
```

```python
logr =LogisticRegression()
logr.fit(fs,target_vector)
```

```python
observation=[[1.4,2.3,5.0,11,12,13,14,15,4,5,7,6,7,13]]
```

```python
prediction=logr.predict(observation)
print(prediction)
```

```python
logr.classes_
```

```python
logr.score(fs,target_vector)
```

```python
logr.predict_proba(observation)[0][0]
```

```python
logr.predict_proba(observation)[0][1]
```

random forest

```python
from sklearn.ensemble import RandomForestClassifier
from sklearn.tree import plot_tree
```

```python
x=df2.drop('station',axis=1)
y=df2['station']
```

```python
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.70)
```

```python
rfc=RandomForestClassifier()
rfc.fit(x_train,y_train)
```

```python
parameters = {'max_depth':[1,2,3,4,5],
              'min_samples_leaf':[5,10,15,20,25],
              'n_estimators':[10,20,30,40,50]}
```

```python
from sklearn.model_selection import GridSearchCV
```

```python
grid_search = GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring='a
grid_search.fit(x_train,y_train)
```

```python
grid_search.best_score_
```

```python
rfc_best =grid_search.best_estimator_
```

```python
py.figure(figsize=(80,50))
plot_tree(rfc_best.estimators_[5],feature_names=x.columns,filled=True)
```

conclusion

The bestfit model is Logistic Regression with score of 0.9102362855254461

```python

```